

Lab 2 BDAT 1008

Loading Data with Schema

Nikhilesh Ramabhadraiah Vashisht
200512808

Create a schema for the data frame using an object of class StructType	2
Question 1	3
Question 2	5
Question 3	6
Question 4	7
Question 5	8

Create a schema for the data frame using an object of class StructType

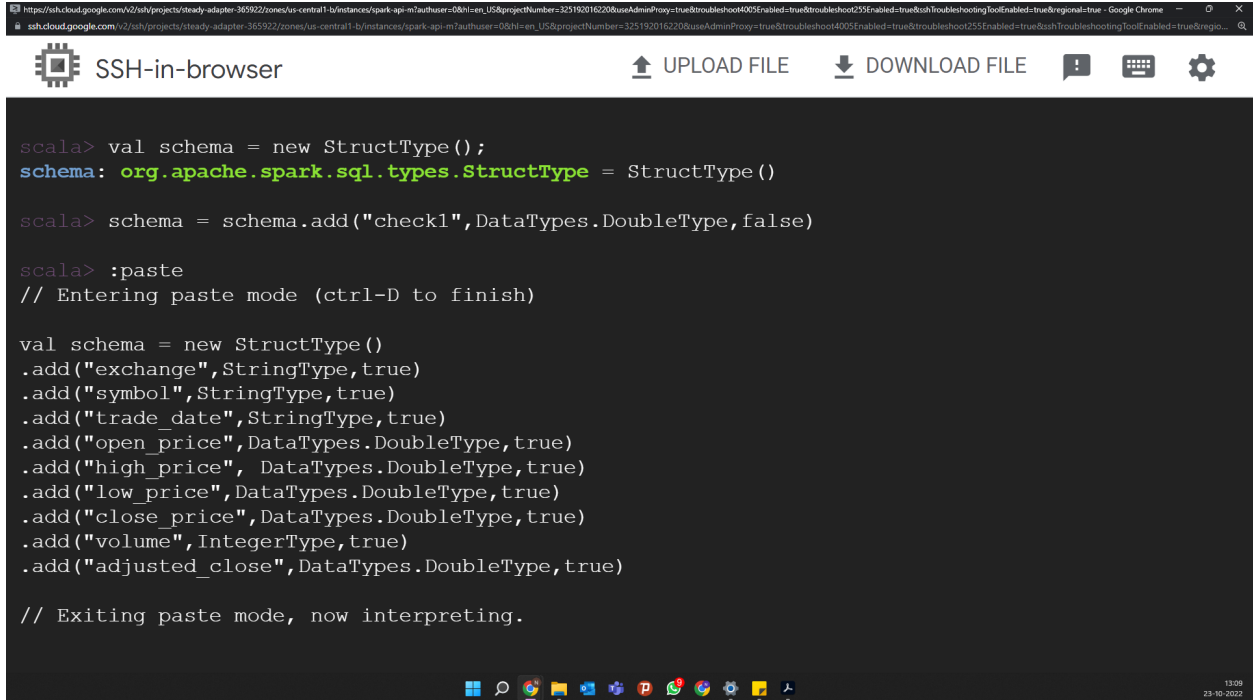
Importing the class:

```
import org.apache.spark.sql.Row;
import org.apache.spark.sql.types.StructField;
import org.apache.spark.sql.types.StructType;
import org.apache.spark.sql.types.StringType;
import org.apache.spark.sql.types.DataTypes._;
```

Specifying schema:

```
val schema = new StructType()
.add("exchange", StringType, true)
.add("symbol", StringType, true)
.add("trade_date", StringType, true)
.add("open_price", DataTypes.DoubleType, true)
.add("high_price", DataTypes.DoubleType, true)
.add("low_price", DataTypes.DoubleType, true)
.add("close_price", DataTypes.DoubleType, true)
.add("volume", IntegerType, true)
.add("adjusted_close", DataTypes.DoubleType, true)
```

BDAT 1008 – Data Collection and Curation - Lab 2



```
scala> val schema = new StructType();
schema: org.apache.spark.sql.types.StructType = StructType()

scala> schema = schema.add("check1", DataTypes.DoubleType, false)

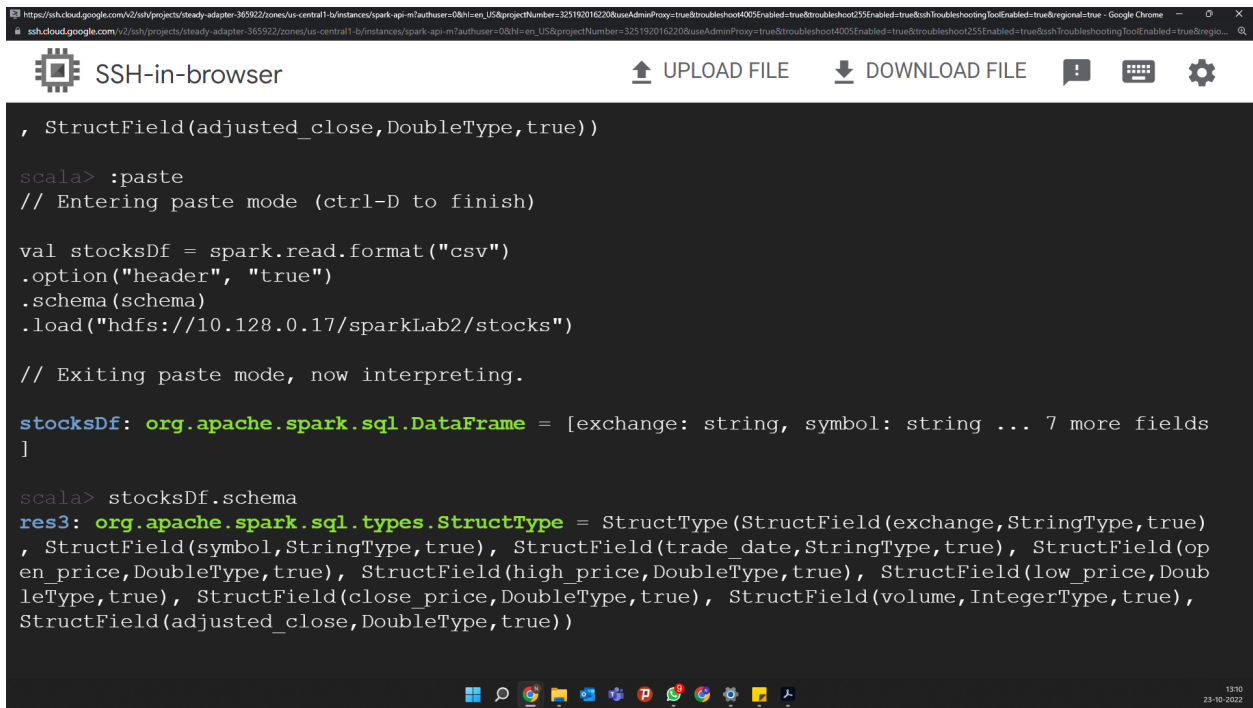
scala> :paste
// Entering paste mode (ctrl-D to finish)

val schema = new StructType()
.add("exchange", StringType, true)
.add("symbol", StringType, true)
.add("trade_date", StringType, true)
.add("open_price", DataTypes.DoubleType, true)
.add("high_price", DataTypes.DoubleType, true)
.add("low_price", DataTypes.DoubleType, true)
.add("close_price", DataTypes.DoubleType, true)
.add("volume", IntegerType, true)
.add("adjusted_close", DataTypes.DoubleType, true)

// Exiting paste mode, now interpreting.
```

Loading the stocks data

```
val stocksDf = spark.read.format("csv")
.option("header", "true")
.schema(schema)
.load("hdfs://10.128.0.17/sparkLab2/stocks")
```



```
, StructField(adjusted_close, DoubleType, true))

scala> :paste
// Entering paste mode (ctrl-D to finish)

val stocksDf = spark.read.format("csv")
.option("header", "true")
.schema(schema)
.load("hdfs://10.128.0.17/sparkLab2/stocks")

// Exiting paste mode, now interpreting.

stocksDf: org.apache.spark.sql.DataFrame = [exchange: string, symbol: string ... 7 more fields]

scala> stocksDf.schema
res3: org.apache.spark.sql.types.StructType = StructType(StructField(exchange,StringType,true),
 StructField(symbol,StringType,true), StructField(trade_date,StringType,true), StructField(open_price,DoubleType,true),
 StructField(high_price,DoubleType,true), StructField(low_price,DoubleType,true), StructField(close_price,DoubleType,true),
 StructField(volume,IntegerType,true), StructField(adjusted_close,DoubleType,true))
```

Question 1

Write a command to find the stocks with average daily volume larger than 1 million shares

val Q1 =

```
stocksDf.select(col("symbol"),col("trade_date"),col("volume")).groupBy(col("symbol")).agg(avg(col("volume"))).filter(avg(col("volume")) > 1000000).show()
```

```
Q1: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [symbol: string, avg(volume): double]

scala> val Q1 = stocksDf.groupBy(col("symbol")).agg(avg(col("volume"))).filter(avg(col("volume")) > 1000000).show()
+-----+
|symbol|      avg(volume)|
+-----+
|BRB|1145559.3268742596|
|BVX|3249145.6251234445|
|BLJ|1162483.5794447726|
|BRE|2752046.230274693|
|BJV|3434707.132243685|
|BVI|3796327.433349109|
|ZVX|3249145.6251234445|
|ZY|1581592.5819723227|
|BAJ|1585092.4|
|BRL|1429770.3473266785|
|ZUY|1272606.6886071048|
|BU|1970861.1921806168|
|BLB|1024276.014976874|
|BRC|5760770.4307449125|
|BX|3570221.4695752007|
|BUY|1272606.6886071048|
|ZX|3570221.4695752007|
|QO7|1825173.7040358745|
|IAJ|1585092.4|
|HVH|1161540.341160833|
+-----+
only showing top 20 rows

Q1: Unit = ()

scala> 
```

Question 2

Write a Scala DataFrame query to find the top 3 stocks by volume for the year 2004.

```
val Q2 =
stocksDf.select(col("exchange"),col("symbol"),col("trade_date"),col("volume")).filter(col("trade_date").contains("2004")).orderBy(col("volume").desc).show(3)
```

The screenshot shows a web browser window with an SSH-in-browser interface. The address bar shows a URL from Google Cloud. The interface includes an 'SSH-in-browser' header, 'UPLOAD FILE' and 'DOWNLOAD FILE' buttons, and a settings icon. The main area displays a Scala REPL session. The first query filters for the year 2004 and orders by volume, showing three rows of data. The second query filters for the year 2004 and orders by volume, showing three rows of data. The session ends with a prompt for the next command.

```
at org.apache.spark.sql.functions$.typedLit(functions.scala:131)
at org.apache.spark.sql.functions$.lit(functions.scala:114)
at org.apache.spark.sql.Column.apply(Column.scala:228)
... 50 elided

scala> val Q2 = stocksDf.select(col("trade_date"),col("volume")).filter(col("trade_date").contains("2004")).orderBy(col("volume").desc).show(3)
+-----+-----+
|trade_date| volume|
+-----+-----+
|2004-09-30|145015500|
|2004-09-30|145015500|
|2004-09-30|145015500|
+-----+-----+
only showing top 3 rows

Q2: Unit = ()

scala> val Q2 = stocksDf.select(col("exchange"),col("symbol"),col("trade_date"),col("volume")).filter(col("trade_date").contains("2004")).orderBy(col("volume").desc).show(3)
+-----+-----+
|exchange|symbol|trade_date| volume|
+-----+-----+
| ABCSE| GRC|2004-09-30|145015500|
| ABCSE| BRC|2004-09-30|145015500|
| ABCSE| IRC|2004-09-30|145015500|
+-----+-----+
only showing top 3 rows

Q2: Unit = ()

scala>
```

Question 3

Write a Scala DataFrame query to find the top 3 stocks by volume and whose symbol start with the first letter of your name (example for Saber, it is symbols starting with “S”).

```
val Q3 =
stocksDf.select(col("exchange"),col("symbol"),col("volume")).filter(col("symbol").like("N%")).orde
rBy(col("volume").desc).show(3)
```

The screenshot shows a Google Cloud Shell terminal window with the following content:

```
scala> val Q3_test = stocksDf.filter(col("symbol").rlike("N%")).show()
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|exchange|symbol|trade_date|open_price|high_price|low_price|close_price|volume|adjusted_close|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+

Q3_test: Unit = ()

scala> val Q3 = stocksDf.select(col("exchange"),col("symbol"),col("volume")).filter(col("symbol").like("N%")).orderBy(col("volume").desc).show(3)
+-----+-----+-----+
|exchange|symbol|volume|
+-----+-----+-----+
|
+-----+-----+-----+

Q3: Unit = ()

scala>
```

There were no symbols starting with the first letter of my name (N)

Question 4

Write a Scala DataFrame to find all the stocks symbols whose closing price is larger than your age.

```
val Q4 =
stocksDf.select(col("exchange"),col("symbol"),col("close_price")).filter(col("close_price")>24).show()

```

The screenshot shows a Google Cloud Shell terminal window with the following content:

```

https://ssh.cloud.google.com/v2/ssh/projects/steady-adaptor-365922/zones/us-central1-b/instances/spark-api-m?authuser=0&hl=en_US&projectNumber=325192016220&useAdminProxy=true&troubleshoot4005Enabled=true&troubleshoot255Enabled=true&sshTroubleshootingToolEnabled=true&regional=...
ssh.cloud.google.com/v2/ssh/projects/steady-adaptor-365922/zones/us-central1-b/instances/spark-api-m?authuser=0&hl=en_US&projectNumber=325192016220&useAdminProxy=true&troubleshoot4005Enabled=true&troubleshoot255Enabled=true&sshTroubleshootingToolEnabled=true&regio...
SSH-in-browser
[+] UPLOAD FILE  [-] DOWNLOAD FILE  [?] [?] [?]
+-----+
| ABCSE | B7B | 36.561 |
| ABCSE | B7B | 36.251 |
| ABCSE | B7B | 36.241 |
| ABCSE | B7B | 36.041 |
| ABCSE | B7B | 36.221 |
| ABCSE | B7B | 36.211 |
| ABCSE | B7B | 36.031 |
| ABCSE | B7B | 35.941 |
| ABCSE | B7B | 36.181 |
| ABCSE | B7B | 36.451 |
| ABCSE | B7B | 36.51 |
| ABCSE | B7B | 35.491 |
| ABCSE | B7B | 35.711 |
| ABCSE | B7B | 36.091 |
| ABCSE | B7B | 35.831 |
| ABCSE | B7B | 35.981 |
| ABCSE | B7B | 36.051 |
| ABCSE | B7B | 36.571 |
| ABCSE | B7B | 36.891 |
+-----+
only showing top 100 rows
Q4: Unit = ()
scala> val Q4 = stocksDf.select(col("exchange"),col("symbol"),col("close_price")).filter(col("close_price")>24).show()
+-----+
|exchange|symbol|close_price|
+-----+
| ABCSE | B7B | 38.681 |
| ABCSE | B7B | 39.041 |
| ABCSE | B7B | 38.241 |
| ABCSE | B7B | 38.321 |
| ABCSE | B7B | 38.511 |
| ABCSE | B7B | 38.251 |
| ABCSE | B7B | 38.221 |
| ABCSE | B7B | 38.341 |
| ABCSE | B7B | 38.581 |
| ABCSE | B7B | 38.081 |
| ABCSE | B7B | 38.111 |
| ABCSE | B7B | 37.491 |
| ABCSE | B7B | 38.011 |
| ABCSE | B7B | 38.311 |
| ABCSE | B7B | 38.81 |
| ABCSE | B7B | 38.781 |
| ABCSE | B7B | 39.011 |
| ABCSE | B7B | 39.351 |
| ABCSE | B7B | 38.821 |
| ABCSE | B7B | 38.971 |
+-----+
only showing top 20 rows
Q4: Unit = ()
scala>

```


Question 5

Write a Scala DataFrame to find the top 10 stocks with largest intraday price change (difference between high and low price during a trading day) and also display the amount of the change.

```
val Q5 =
stocksDf.select(col("exchange"),col("symbol"),col("high_price"),col("low_price"),(col("high_price")
)-col("low_price")).alias("price_change_in_a_day")).orderBy(col("price_change_in_a_day").desc
).show(10)
```

```

only showing top 20 rows
Q4: Unit = ()

scala> val Q5 = stocksDf.select(col("exchange"),col("symbol"),col("high_price"),col("low_price"),(col("high_price")-col("low_price")).alias("price_change_in_a_day")).show()
-----
[exchange|symbol|high_price|low_price|price_change_in_a_day|
-----
| ABCSE| B7J| 8.71| 8.57| 0.13999999999999999|
| ABCSE| B7J| 8.71| 8.31| 0.400000000000000036|
| ABCSE| B7J| 8.88| 8.59| 0.29000000000000009|
| ABCSE| B7J| 8.92| 8.8| 0.119999999999999922|
| ABCSE| B7J| 8.9| 8.73| 0.16999999999999993|
| ABCSE| B7J| 8.77| 8.66| 0.109999999999999943|
| ABCSE| B7J| 8.81| 8.56| 0.25|
| ABCSE| B7J| 8.9| 8.61| 0.30000000000000007|
| ABCSE| B7J| 8.87| 8.68| 0.18999999999999995|
| ABCSE| B7J| 8.92| 8.71| 0.20999999999999998|
| ABCSE| B7J| 9.0| 8.73| 0.26999999999999996|
| ABCSE| B7J| 9.04| 8.94| 0.09999999999999994|
| ABCSE| B7J| 9.13| 8.98| 0.150000000000000036|
| ABCSE| B7J| 9.09| 8.93| 0.160000000000000014|
| ABCSE| B7J| 9.04| 8.91| 0.12999999999999998|
| ABCSE| B7J| 9.0| 8.9| 0.09999999999999994|
| ABCSE| B7J| 9.03| 8.93| 0.09999999999999994|
| ABCSE| B7J| 9.0| 8.92| 0.08000000000000007|
| ABCSE| B7J| 8.95| 8.86| 0.08999999999999996|
| ABCSE| B7J| 9.0| 8.95| 0.050000000000000071|
-----

only showing top 20 rows
Q5: Unit = ()

scala> val Q5 = stocksDf.select(col("exchange"),col("symbol"),col("high_price"),col("low_price"),(col("high_price")-col("low_price")).alias("price_change_in_a_day")).orderBy(col("price_change_in_a_day").desc).show(10)
-----
[exchange|symbol|high_price|low_price|price_change_in_a_day|
-----
| ABCSE| BBR| 583.51| 475.17| 108.33999999999997|
| ABCSE| BBR| 583.51| 475.17| 108.33999999999997|
| ABCSE| BBR| 583.51| 475.17| 108.33999999999997|
| ABCSE| BBR| 583.51| 475.17| 108.33999999999997|
| ABCSE| BBR| 583.51| 475.17| 108.33999999999997|
| ABCSE| CCL| 480.0| 380.1| 99.89999999999998|
| ABCSE| CCL| 480.0| 380.1| 99.89999999999998|
| ABCSE| CCL| 480.0| 380.1| 99.89999999999998|
| ABCSE| BCL| 480.0| 380.1| 99.89999999999998|
| ABCSE| BBR| 421.0| 338.66| 82.33999999999997|
-----

only showing top 10 rows
Q5: Unit = ()

scala>

```